

SD-TSIA204

Statistics : linear models

Ekhine Irurozki
Télécom Paris

Grades : SD-TSIA204

- Practical 2 : **20% final grade**
 - Problem statement available on Wednesday 19/01/2021, deadline 04/02/2021 at 00 :01.
 - Single accepted file : **IPython Notebook**
- Final exam : **80% note finale**
 - Date : 02/02/2021
 - Format : Quiz + exercises

Rem:Quiz questions samples are available on the course website (cf. Liste de questions section)

BEWARE : your practical should be personal not copy pasted from your neighbor !!!

Practical notation

Practicals are graded on a scale from 0 to **20**, as follows

- ▶ scientific quality of answer **15** pts
- ▶ language/writing quality of answer (spelling, etc.) **2** pts
- ▶ indentation, PEP8 Style, useful comments in code, no/few warnings **2** pts
- ▶ no bug **1** pt (at least
https://github.com/agramfort/check_notebook)
- ▶ one single **.ipynb** file expected, submitted on the “Site pédagogique” of the course ; emailed work will receive a zero score and will not be graded

Late : **no Late work** work will be accepted, unless official reason accepted by Télécom ParisTech’s administration ; late work will receive a zero score and will not be graded

Prerequisites

- ▶ **Probability** basis : probability, expectation, law of large number, Gaussian distribution, central limit theorem.
Books : Foata et Fuchs (1996) (in French) or Murphy (2012, ch.1 and 2)
- ▶ **Optimisation** basis : (differential) calculus, convexity, first order conditions, gradient descent, Newton method
Lecture : Boyd et Vandenberghe (2004), Bertsekas (1999)
- ▶ **(bi-)linear algebra** basis : vector space, norms, inner product, matrices, determinants, diagonalization
Lecture : Horn et Johnson (1994)
- ▶ **Numerical linear algebra** : linear system resolution, Gaussian elimination, matrix factorization, conditioning, etc.
Lecture : Golub et VanLoan (2013), [link](#) par L. Vandenberghe
- ▶ **Numerical linear algebra** : Introduction to Applied Linear Algebra – Vectors, Matrices, and Least Squares Lecture :
Boyd et Vandenberghe (2018), [link](#) par Stephen Boyd and Lieven Vandenberghe

Algorithmic aspects : some advice

Python installation : use **Conda** / **Anaconda**

Rem: you are on your own for this (or use the school machines)

Recommended tools : **Jupyter** / **IPython Notebook** (mandatory for your practical) **IPython** with a text editor e.g., **Atom**, **Sublime Text**, **Visual Studio Code**, etc., for larger projects

- ▶ **Python, Scipy, Numpy** :

http://perso.telecom-paristech.fr/~gramfort/liesse_python/

- ▶ **Pandas** : <http://pandas.pydata.org/>

- ▶ **scikit-learn** : <http://scikit-learn.org/stable/>

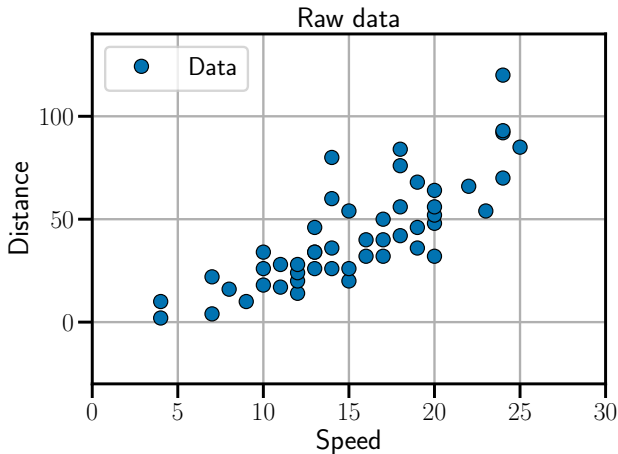
Rem: for practicals, bring your own machine if you prefer but install your Python environment upfront

General advice

- ▶ Use version control system for your work : **Git** (e.g., **Bitbucket**, **Github**, etc.)
- ▶ Use clean way of writing code/ presenting your code
Example : **PEP8** for Python (use for instance **AutoPEP8**, <https://github.com/kenko000/jupyter-autopep8>)
- ▶ Learn from good examples :
<https://github.com/scikit-learn/>,
<http://jakevdp.github.io/>, etc.

A 2D starting example

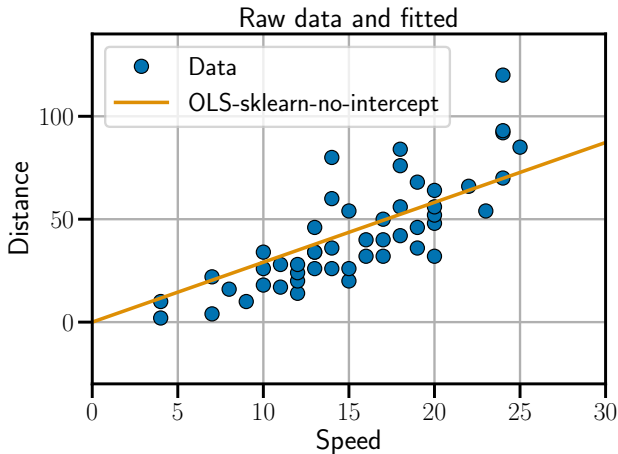
Example : braking distance for cars as a function of speed
($n = 50$ measurements)



Dataset *cars* : <https://stat.ethz.ch/R-manual/R-devel/library/datasets/html/cars.html>
<https://stat.ethz.ch/R-manual/R-devel/library/datasets/html/cars.html>

A 2D starting example

Example : braking distance for cars as a function of speed
($n = 50$ measurements)



Dataset *cars* : <https://stat.ethz.ch/R-manual/R-devel/library/datasets/html/cars.html>
<https://stat.ethz.ch/R-manual/R-devel/library/datasets/html/cars.html>

Python command

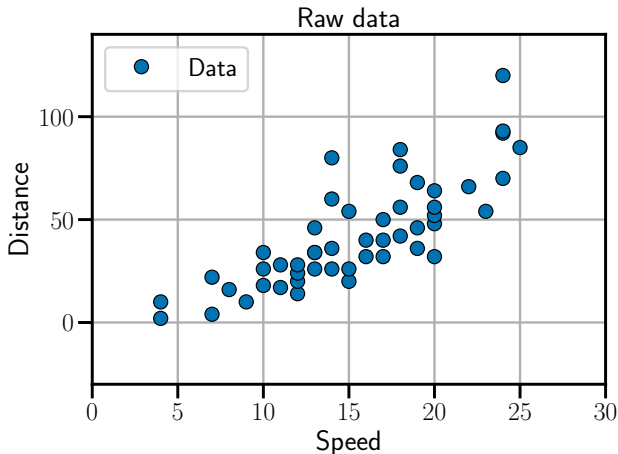
```
import pandas as pd
import matplotlib.pyplot as plt
import sklearn.linear_model as lm
# Load data
url = 'cars.csv'
dat = pd.read_csv(url)
y = dat['dist']
X = dat[['speed']] # sklearn needs X to have 2 dim.

skl_linmod = lm.LinearRegression(fit_intercept=False)
skl_linmod.fit(X, y) # Fit regression model

fig = plt.figure(figsize=(8, 6))
plt.plot(X, y, 'o', label="Data")
plt.plot(X, skl_linmod.predict(X),
         label="OLS-sklearn-no-intercept")
plt.legend(loc='upper left')
plt.show()
```

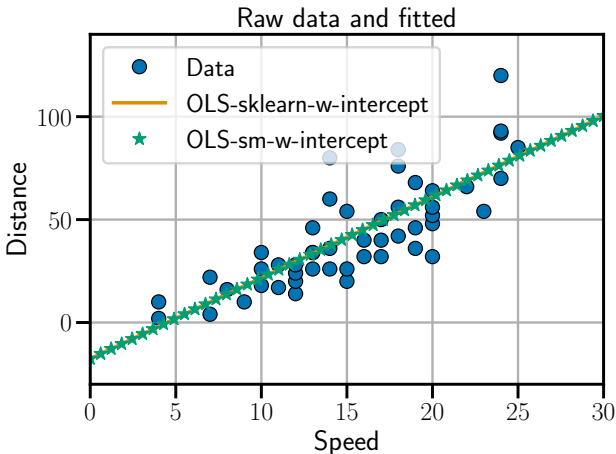
A 2D starting example : with an intercept

Example : braking distance for cars as a function of speed ($n = 50$ measurements)



A 2D starting example : with an intercept

Example : braking distance for cars as a function of speed ($n = 50$ measurements)



Python commands : with intercept

```
import statsmodels.api as sm

# data, fitted, etc
y = dat['dist']
X = dat[['speed']]
X = sm.add_constant(X)
results = sm.OLS(y,X).fit()

# plot
fig, ax = plt.subplots(figsize=(8,6))
ax.plot(X['speed'], y, 'o', label="data")
ax.plot(X['speed'], results.fittedvalues,
        linewidth=3, label="OLS-sm-w-intercept")
ax.legend(loc='best')
```

Alternative : use `lm.LinearRegression(fit_intercept=True)`

Notation interpretation

- ▶ $n = 50$
- ▶ $p = 1$
- ▶ y_i : braking time for i -th car
- ▶ x_i : speed of i -th car
- ▶ y : the observation is the car's braking time
- ▶ x : the feature/covariate is the car's speed

Linear model / Linear regression hypothesis : assume that braking time is proportional to speed

Exercise: use `describe()` from Pandas to get a rough data summary

Modeling I, the 1D case

Given a sample : (y_i, x_i) , for $i = 1, \dots, n$

Linear model or linear regression hypothesis assume :

$$y_i \approx \theta_0^* + \theta_1^* x_i$$

Model

- ▶ θ_0^* : intercept (unknown)
- ▶ θ_1^* : slope (unknown)

Rem: both parameters are unknown from the statistician

Data

- ▶ y is an **observation** or a variable to explain
- ▶ x is a **feature** or a covariate

Modeling II

Probabilistic model. Let us give a precise meaning to the sign \approx :

$$y_i = \theta_0^* + \theta_1^* x_i + \varepsilon_i,$$

$$\varepsilon_i \stackrel{i.i.d}{\sim} \varepsilon, \text{ for } i = 1, \dots, n$$

$$\mathbb{E}(\varepsilon) = 0$$

where i.i.d. means “independent and identically distributed”

Interpretation : $\varepsilon_i = y_i - \theta_0^* - \theta_1^* x_i$: represent the error between the theoretical model and the observations, represented by random variables ε_i centered (often referred to as **white noise**). Rem: motivation for the random nature of the noise – measurement noise, transmission noise, in-population variability, etc.

Modeling III

$$y_i = \theta_0^\star + \theta_1^\star x_i + \varepsilon_i$$

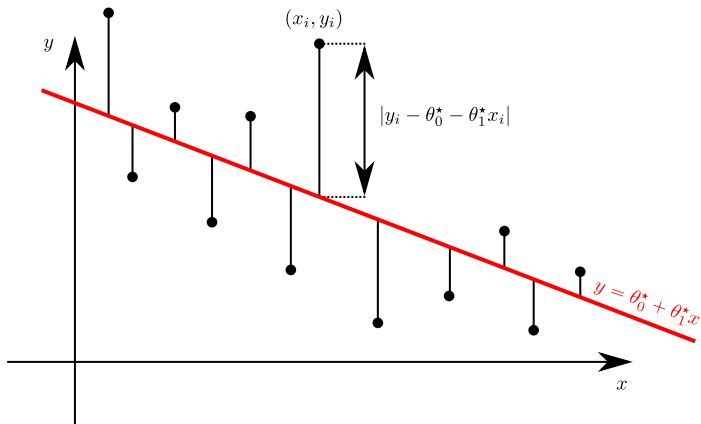
We call

- ▶ **intercept** the scalar θ_0^\star (■ ■ : *ordonnée à l'origine*)
- ▶ **slope** the scalar θ_1^\star (■ ■ : *pente*)

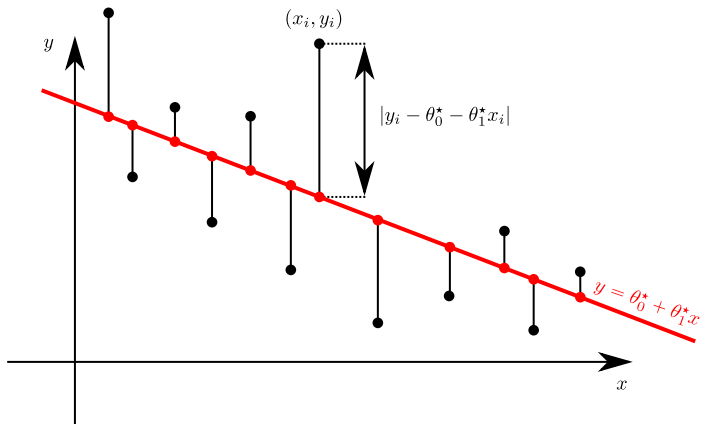
Our **goal** is to estimate θ_0^\star and θ_1^\star (unknown) by $\hat{\theta}_0$ and $\hat{\theta}_1$ relying on observations (y_i, x_i) for $i = 1, \dots, n$

Rem: The “hat” notation is classical in statistics for referring to estimators

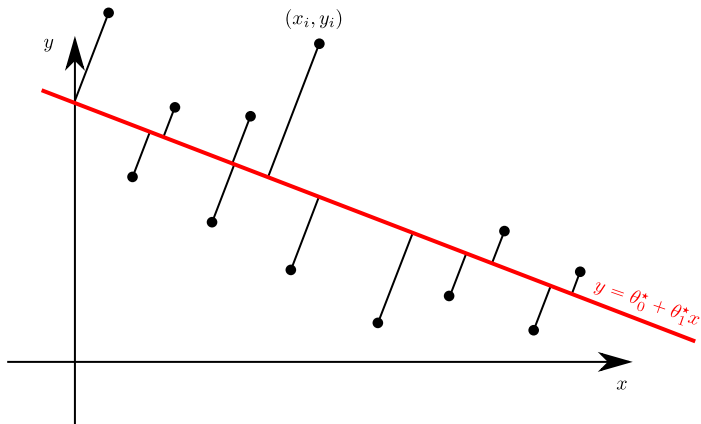
Least squares : visualization



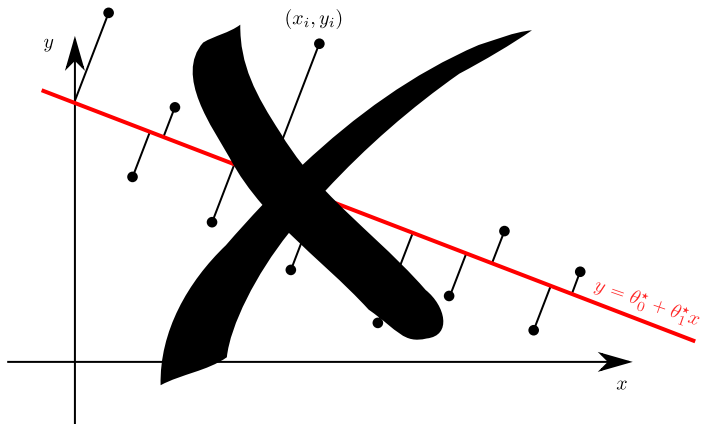
Least squares : visualization



(Total) Least squares : visualization



(Total) Least squares : visualization



Least squares : mathematical formulation

The **least squares** estimator is defined as :

$$(\hat{\theta}_0, \hat{\theta}_1) \in \arg \min_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

- ▶ it is also referred to as “ordinary least squares” (OLS)
- ▶ an original motivation for the squares is computational : first order conditions only require solving a linear system
- ▶ a solution always exists : minimizing a **coercive** continuous function (coercive : $\lim_{\|x\| \rightarrow +\infty} f(x) = +\infty$)

Rem: write « $\in \arg \min$ » as long as you do not know if the solution is unique

Least square authorship (controversial)



FIGURE – Adrien-Marie Legendre and Carl Friedrich Gauss

Historical / robust detour

Definition

The **least absolute deviation** (LAD) estimator reads :

$$(\hat{\theta}_0, \hat{\theta}_1) \in \arg \min_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n |y_i - \theta_0 - \theta_1 x_i|$$

Rem: hard to compute without computer ; requires an optimization solver for non-smooth function (or a Linear Programming solver)

Rem: more robust to outliers (🇫🇷 : *données aberrantes*)

Least absolute deviation authorship

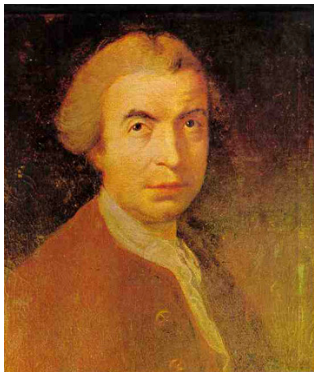
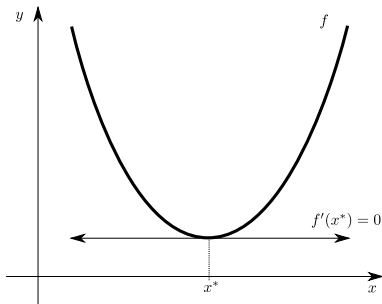


FIGURE – Ruđer Josip Bošković and Pierre-Simon de Laplace

Local minimum : first order condition

Theorem : Fermat's rule

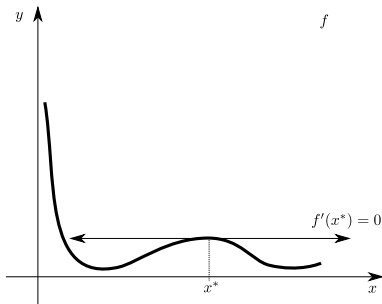
If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , i.e., $\nabla f(x^*) = 0$.



Local minimum : first order condition

Theorem : Fermat's rule

If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , i.e., $\nabla f(x^*) = 0$.

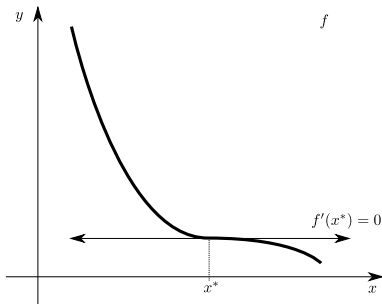


Rem: sufficient condition when f is convex !

Local minimum : first order condition

Theorem : Fermat's rule

If f is differentiable, then at a local minimum x^* the gradient of f vanishes at x^* , i.e., $\nabla f(x^*) = 0$.



Rem: sufficient condition when f is convex !

Back to least squares

$$\hat{\boldsymbol{\theta}} = (\hat{\theta}_0, \hat{\theta}_1) \in \arg \min_{(\theta_0, \theta_1) \in \mathbb{R}^2} \frac{1}{2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

For least squares, minimize the function of two variables :

$$f(\theta_0, \theta_1) = f(\boldsymbol{\theta}) = \frac{1}{2} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

First order condition / Fermat's rule :

$$\begin{cases} \frac{\partial f}{\partial \theta_0}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) = 0 \\ \frac{\partial f}{\partial \theta_1}(\hat{\boldsymbol{\theta}}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) x_i = 0 \end{cases}$$

Exercise: Is f convex? help : a sum of convex functions is convex

The solution is unique ($\Leftrightarrow f$ is strictly convex)

F is quadratic $\implies f$ is convex. Moreover, $\det(\nabla^2 f(\hat{\theta})) > 0 \Leftrightarrow 0$ is not an eigenvalue of $\det(\nabla^2 f(\hat{\theta})) \Leftrightarrow \nabla^2 f(\hat{\theta})$ p.s.d. $\implies f(\hat{\theta})$ strictly convex \Leftrightarrow unique solution.

For

$$\nabla^2 f(\hat{\theta}) = \begin{bmatrix} 1 & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n x_i^2 \end{bmatrix} \quad (1)$$

we have

$$\det(\nabla^2 f(\hat{\theta})) \neq 0 \Leftrightarrow \det(\nabla^2 f(\hat{\theta})/n) \neq 0 \Leftrightarrow \frac{\sum x_i^2}{n} - \left(\frac{\sum x_i}{n}\right)^2 \neq 0 \Leftrightarrow n^{-1} \sum (x_i - \bar{x})^2 \neq 0 \quad \text{i.e., the variance} \neq 0 \quad (2)$$

Conclusion : uniqueness is guaranteed as long as the variance is different from zero, i.e. the values of x_i are not reduced to a single point.

Calculus continued

Usual mean notation : $\bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i$ and $\bar{y}_n = \frac{1}{n} \sum_{i=1}^n y_i$

With that, Fermat's rule states (dividing by n) :

$$\begin{cases} \frac{\partial f}{\partial \theta_0}(\hat{\theta}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) = 0 \\ \frac{\partial f}{\partial \theta_1}(\hat{\theta}) = \sum_{i=1}^n (y_i - \hat{\theta}_0 - \hat{\theta}_1 x_i) x_i = 0 \end{cases}$$

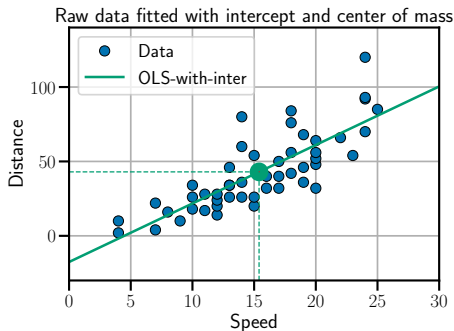
\Leftrightarrow

$$\begin{cases} \hat{\theta}_0 = \bar{y}_n - \hat{\theta}_1 \bar{x}_n & \text{(CNO1)} \\ \hat{\theta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sum_{i=1}^n (x_i - \bar{x}_n)^2} & \text{(CNO2)} \end{cases}$$

Exercise: Prove that (CNO2) holds if and only if $\mathbf{x} = (x_1, \dots, x_n)^\top$ is non constant, i.e., \mathbf{x} is not proportional to $\mathbf{1}_n = (1, \dots, 1)^\top \in \mathbb{R}^n$

Center of gravity and interpretation

$$(\text{CNO1}) \Leftrightarrow (\bar{x}_n, \bar{y}_n) \in \{(x, y) \in \mathbb{R}^2 : y = \hat{\theta}_0 + \hat{\theta}_1 x\}$$



- ▶ $\overline{speed} = 15.4$
- ▶ $\overline{dist} = 42.98$
- ▶ $\hat{\theta}_0 = -17.579095$ intercept (negative!)
- ▶ $\hat{\theta}_1 = 3.932409$ slope

Physical interpretation : the cloud of points' center of gravity belongs to the (estimated) regression line

Vector formulation

Notation : $\mathbf{x} = (x_1, \dots, x_n)^\top$ and $\mathbf{y} = (y_1, \dots, y_n)^\top$

$$(\text{CNO2}) \Leftrightarrow \hat{\theta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sum_{i=1}^n (x_i - \bar{x}_n)^2}$$

$$(\text{CNO2}) \Leftrightarrow \hat{\theta}_1 = \text{corr}_n(\mathbf{x}, \mathbf{y}) \cdot \frac{\sqrt{\text{var}_n(\mathbf{y})}}{\sqrt{\text{var}_n(\mathbf{x})}}$$

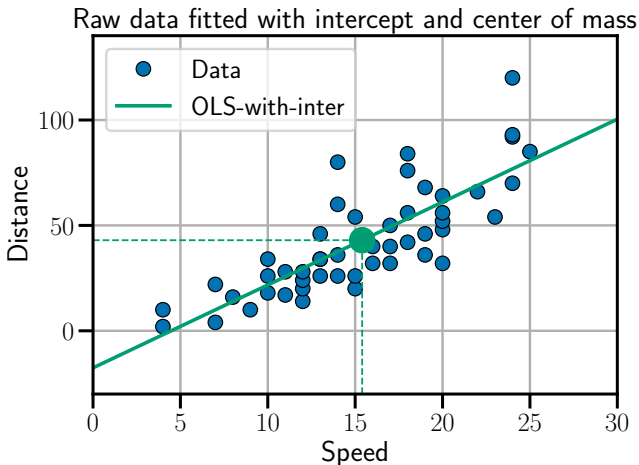
where $\text{corr}_n(\mathbf{x}, \mathbf{y}) = \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x}_n)(y_i - \bar{y}_n)}{\sqrt{\text{var}_n(\mathbf{x})} \sqrt{\text{var}_n(\mathbf{y})}}$

and $\text{var}_n(\mathbf{z}) = \frac{1}{n} \sum_{i=1}^n (z_i - \bar{z}_n)^2$ (for any $\mathbf{z} = (z_1, \dots, z_n)^\top$)

respectively **empirical correlation** **empirical variances**

Back to the *cars* example

Line slope : $\text{corr}_n(\mathbf{x}, \mathbf{y}) \cdot \frac{\sqrt{\text{var}_n(\mathbf{y})}}{\sqrt{\text{var}_n(\mathbf{x})}} = 3.932409$.



Centering

Centered model :

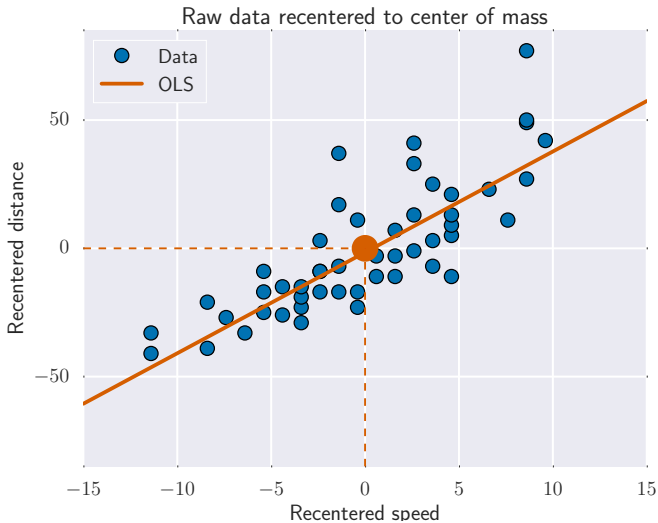
$$\text{Write for any } i = 1, \dots, n : \begin{cases} x'_i = x_i - \bar{x}_n \\ y'_i = y_i - \bar{y}_n \end{cases} \Leftrightarrow \begin{cases} \mathbf{x}' = \mathbf{x} - \bar{x}_n \mathbf{1}_n \\ \mathbf{y}' = \mathbf{y} - \bar{y}_n \mathbf{1}_n \end{cases}$$

and $\mathbf{1}_n = (1, \dots, 1)^\top \in \mathbb{R}^n$, then solving the OLS with $(\mathbf{x}', \mathbf{y}')$ leads to

$$\begin{cases} \hat{\theta}'_0 = 0 \\ \hat{\theta}'_1 = \frac{\frac{1}{n} \sum_{i=1}^n x'_i y'_i}{\frac{1}{n} \sum_{i=1}^n x_i'^2} \end{cases}$$

Rem: equivalent to choosing the cloud of points' center of mass as origin, i.e., $(\bar{x}'_n, \bar{y}'_n) = (0, 0)$

Centering (II)




Centering and interpretation

Consider the coefficient $\hat{\theta}'_1$ ($\hat{\theta}'_0 = 0$) for centered points \mathbf{y}' , \mathbf{x}' , then :

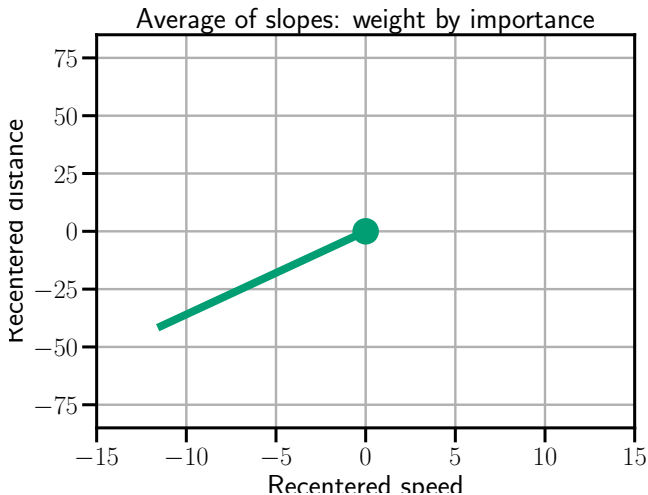
$$\hat{\theta}'_1 \in \arg \min_{\theta_1} \sum_{i=1}^n (y'_i - \theta_1 x'_i)^2 = \arg \min_{\theta_1} \sum_{i=1}^n x'^2_i \left(\frac{y'_i}{x'_i} - \theta_1 \right)^2$$

Interpretation : $\hat{\theta}'_1$ is a weighted average of the slopes $\frac{y'_i}{x'_i}$

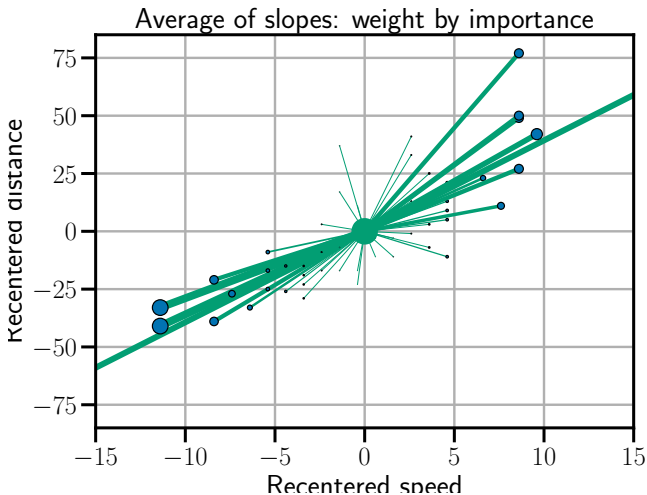
$$\hat{\theta}'_1 = \frac{\sum_{i=1}^n x'^2_i \frac{y'_i}{x'_i}}{\sum_{j=1}^n x'^2_j}$$

Influence of extreme points : weights proportional to x'^2_i ;
connected to the **leverage** ( : *levier*) effect

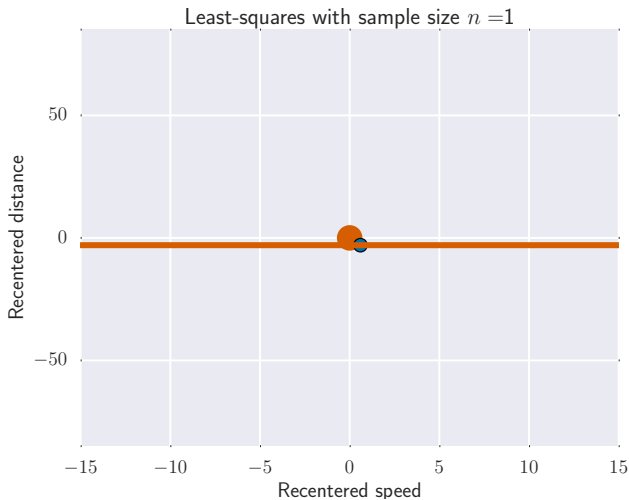
Extreme points – leverage effect



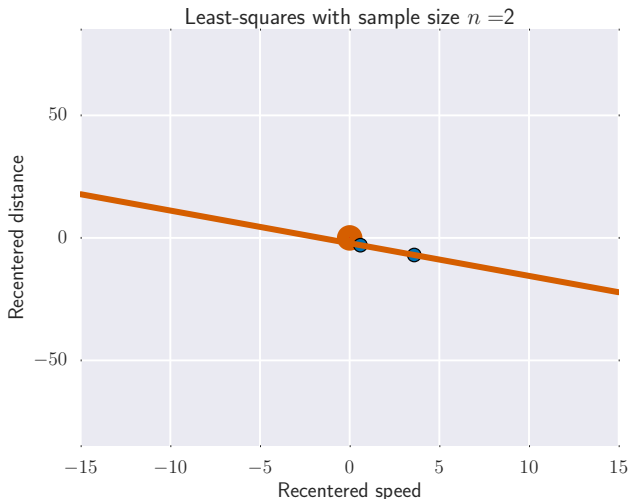
Extreme points – leverage effect



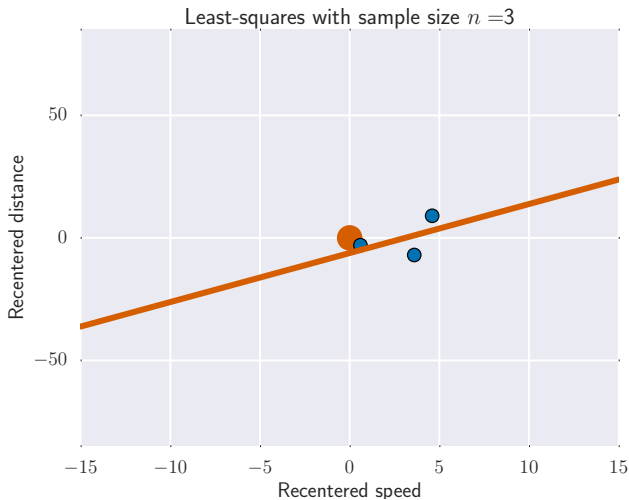
Extreme points – leverage effect (II)



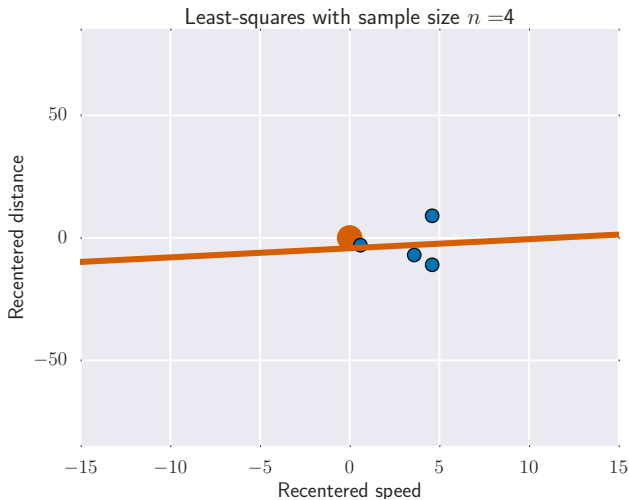
Extreme points – leverage effect (II)



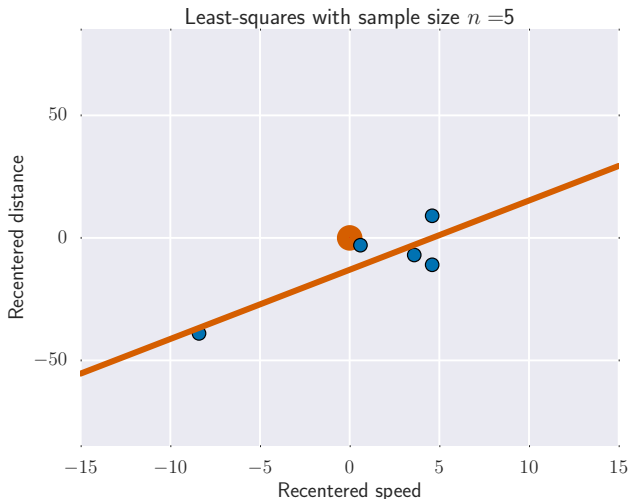
Extreme points – leverage effect (II)



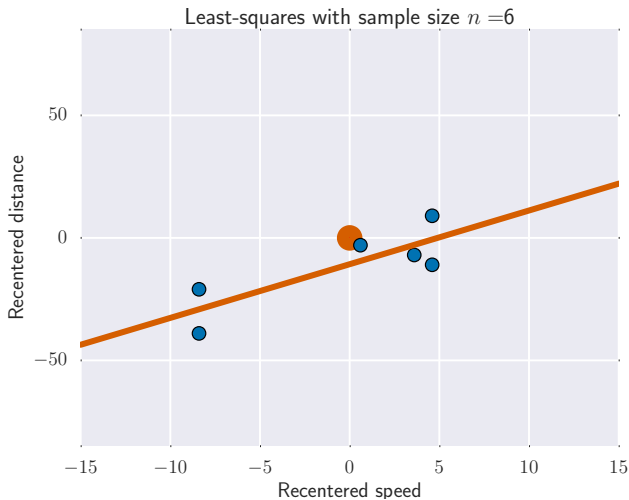
Extreme points – leverage effect (II)



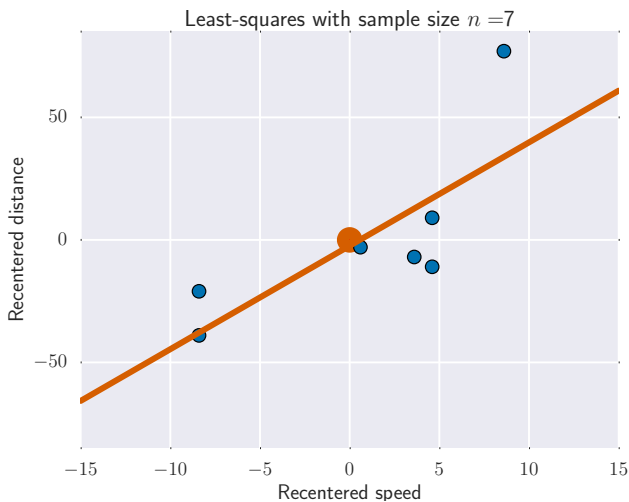
Extreme points – leverage effect (II)



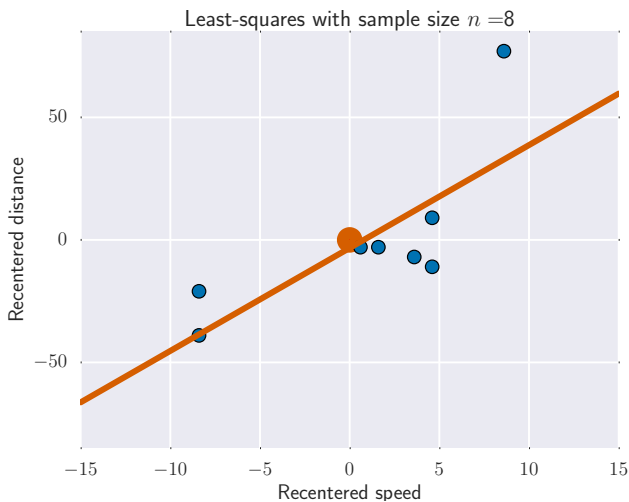
Extreme points – leverage effect (II)



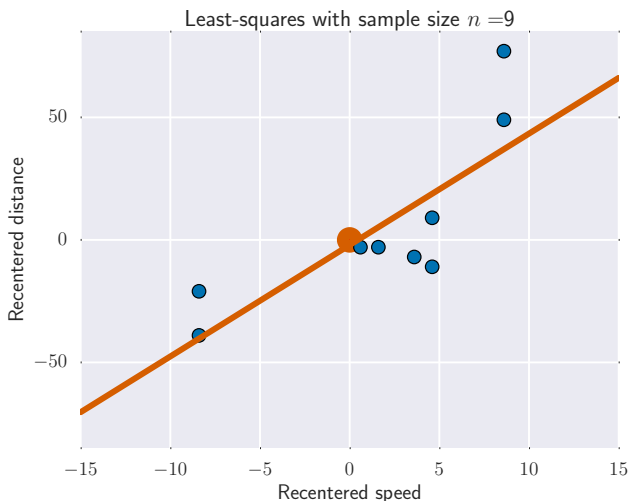
Extreme points – leverage effect (II)



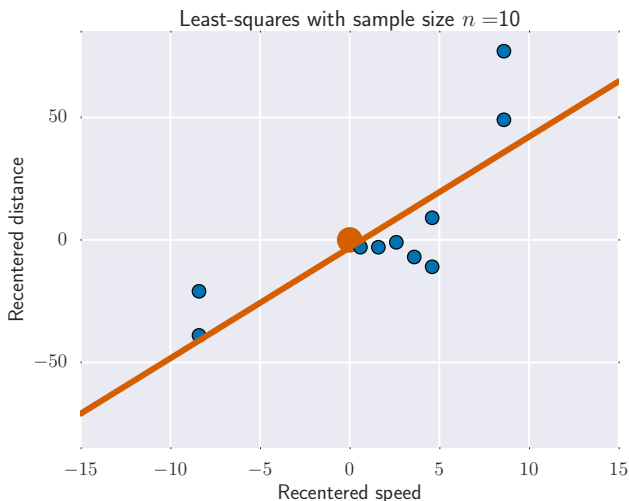
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



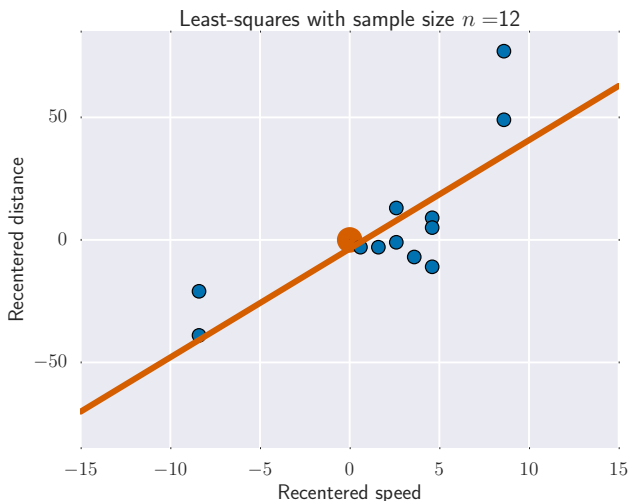
Extreme points – leverage effect (II)



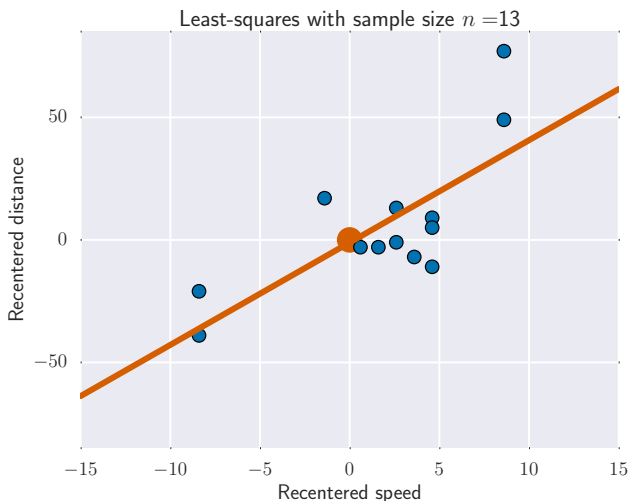
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



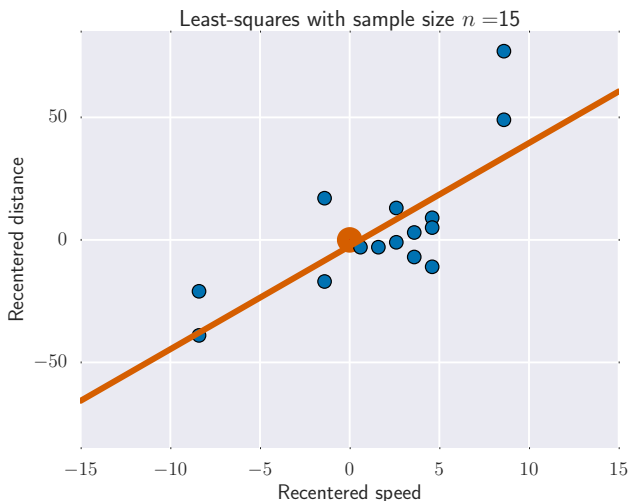
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



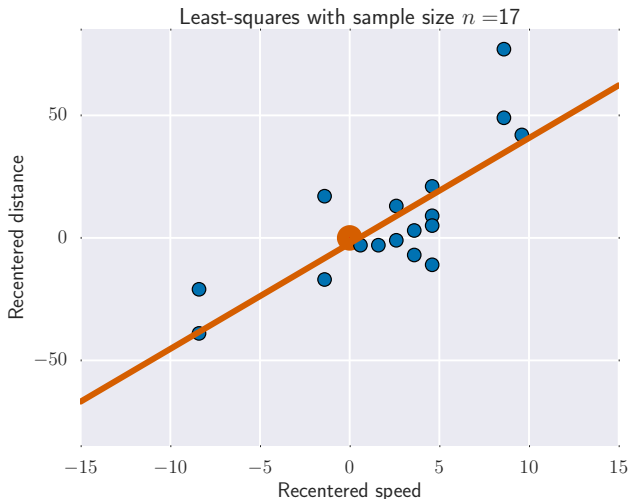
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



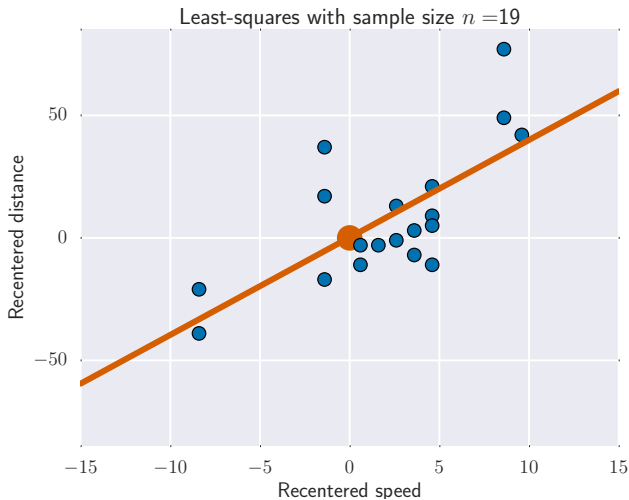
Extreme points – leverage effect (II)



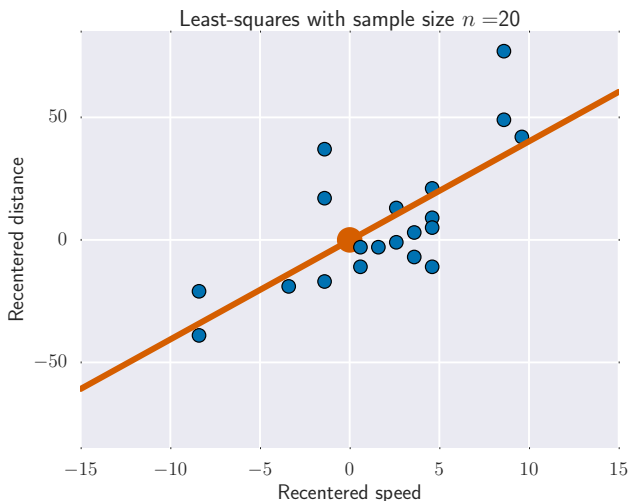
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



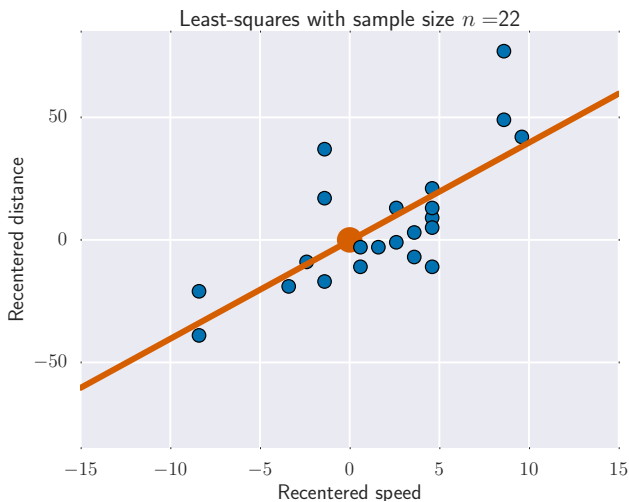
Extreme points – leverage effect (II)



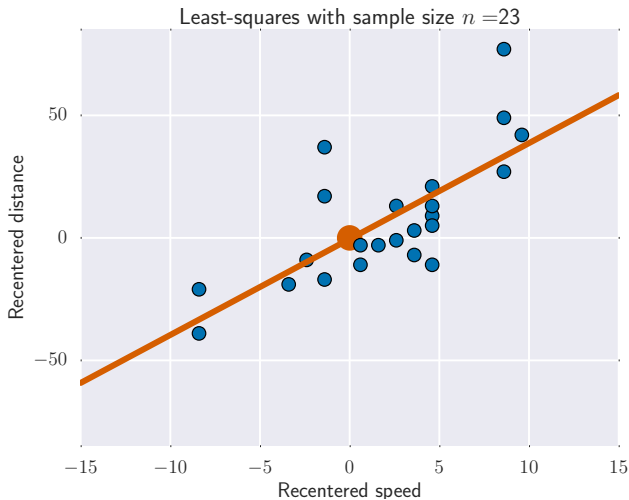
Extreme points – leverage effect (II)



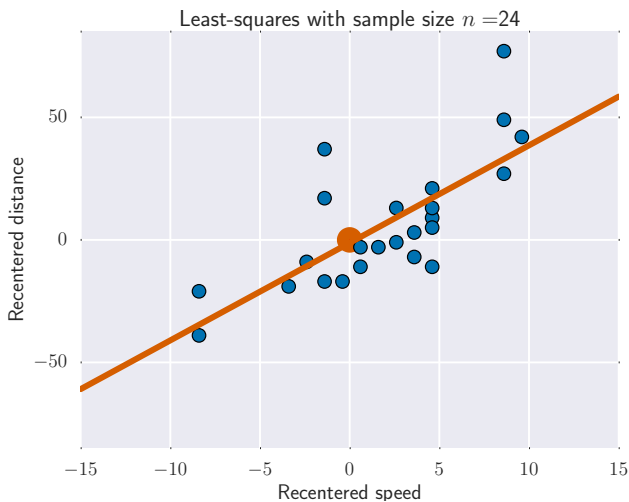
Extreme points – leverage effect (II)



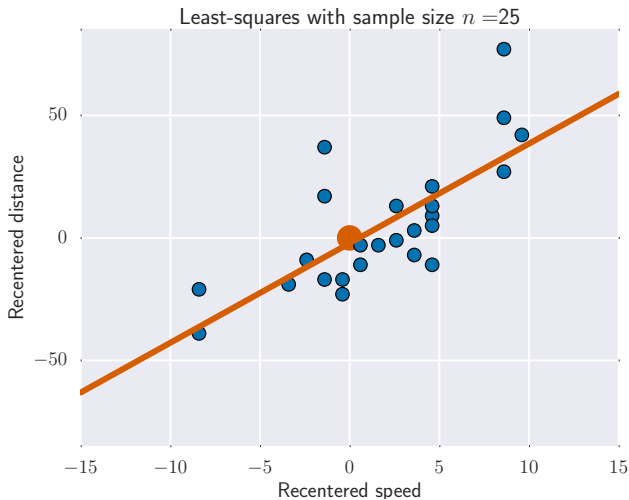
Extreme points – leverage effect (II)



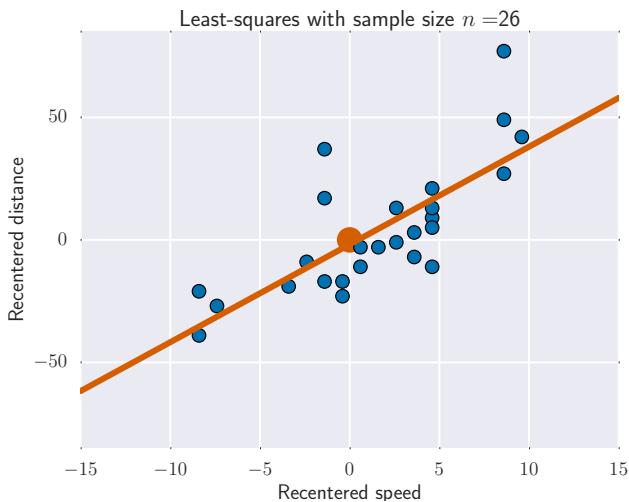
Extreme points – leverage effect (II)



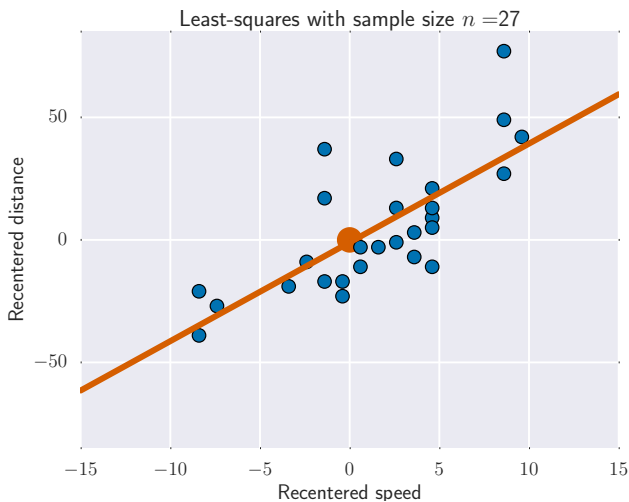
Extreme points – leverage effect (II)



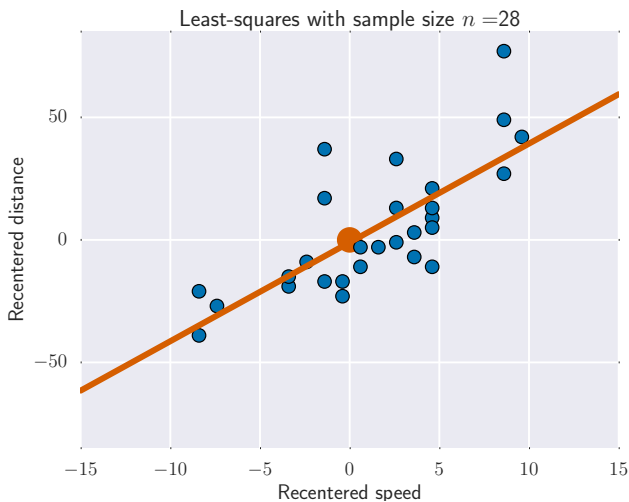
Extreme points – leverage effect (II)



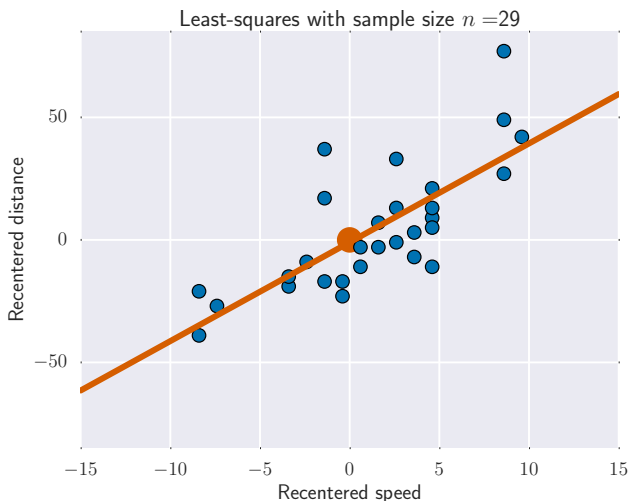
Extreme points – leverage effect (II)



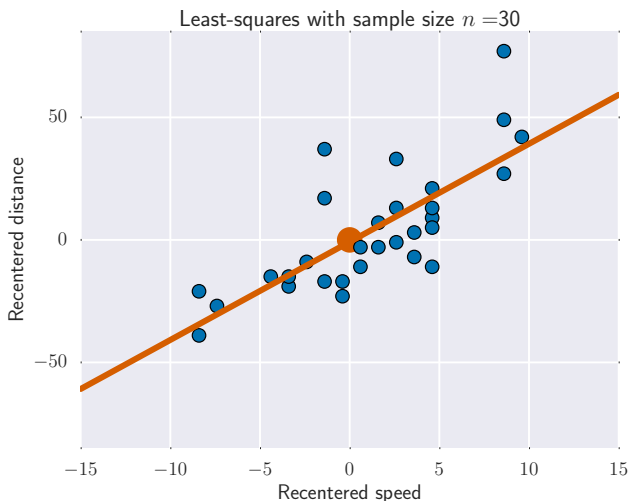
Extreme points – leverage effect (II)



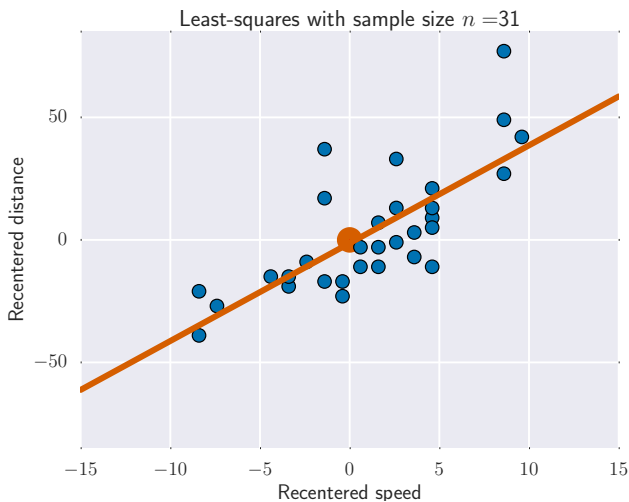
Extreme points – leverage effect (II)



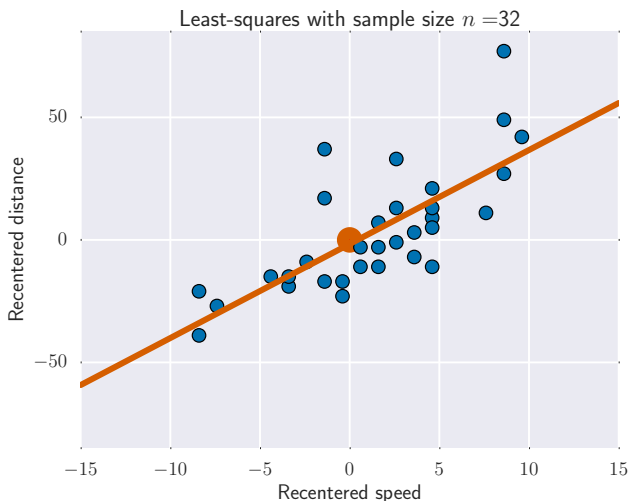
Extreme points – leverage effect (II)



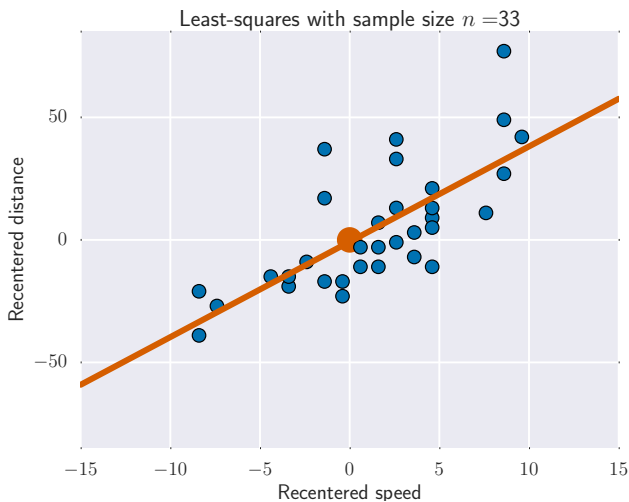
Extreme points – leverage effect (II)



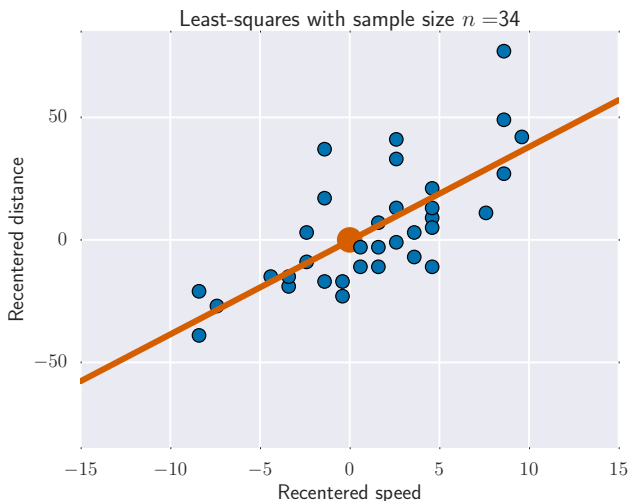
Extreme points – leverage effect (II)



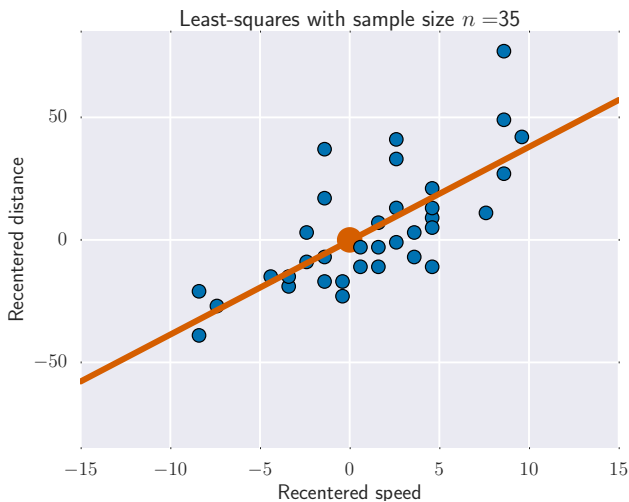
Extreme points – leverage effect (II)



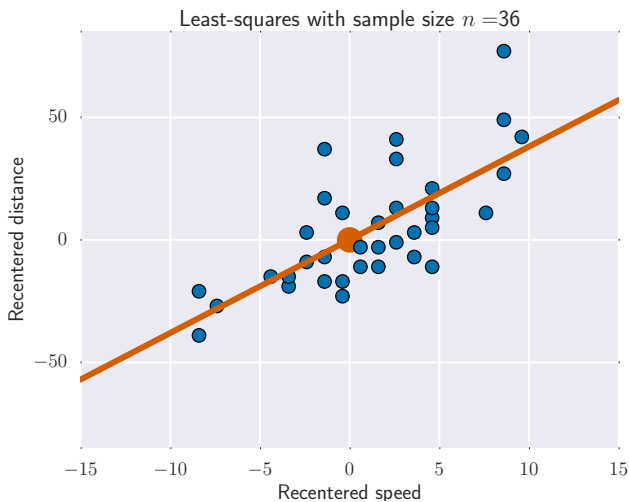
Extreme points – leverage effect (II)



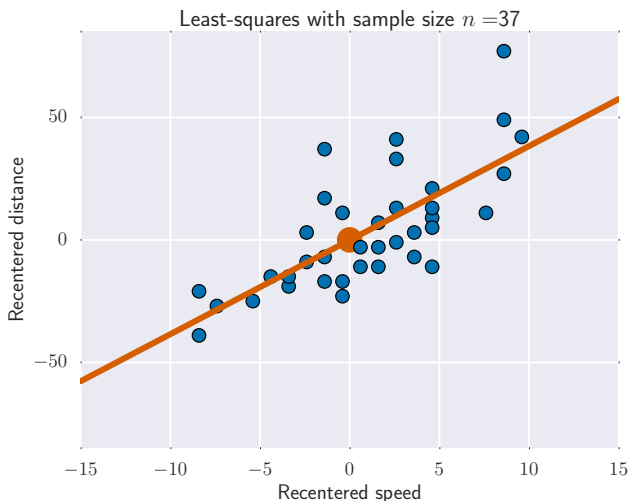
Extreme points – leverage effect (II)



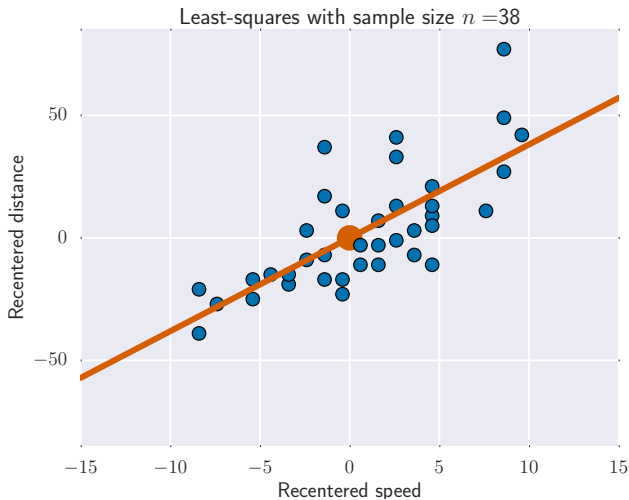
Extreme points – leverage effect (II)



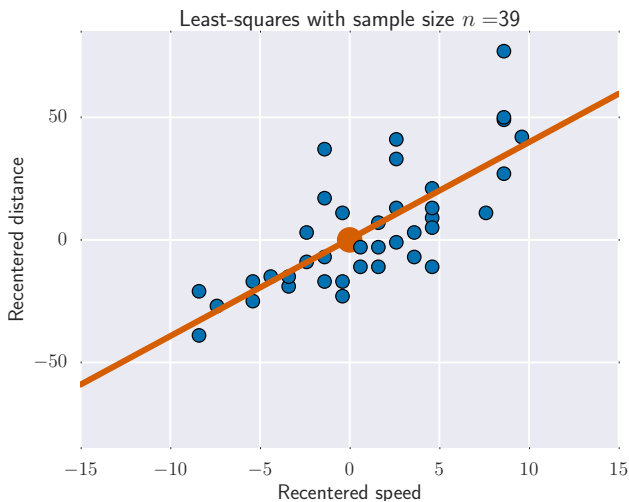
Extreme points – leverage effect (II)



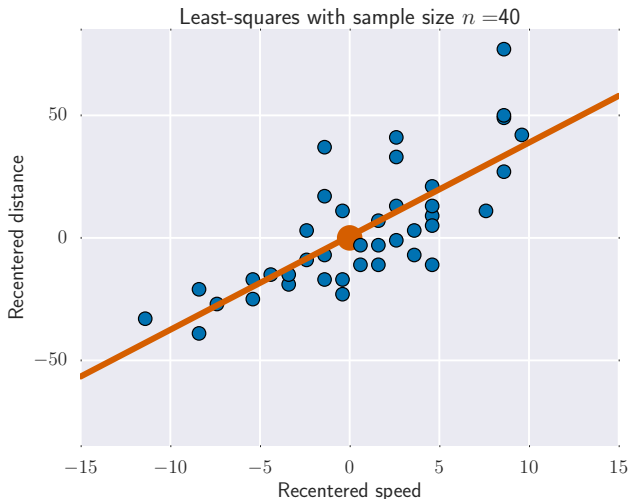
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



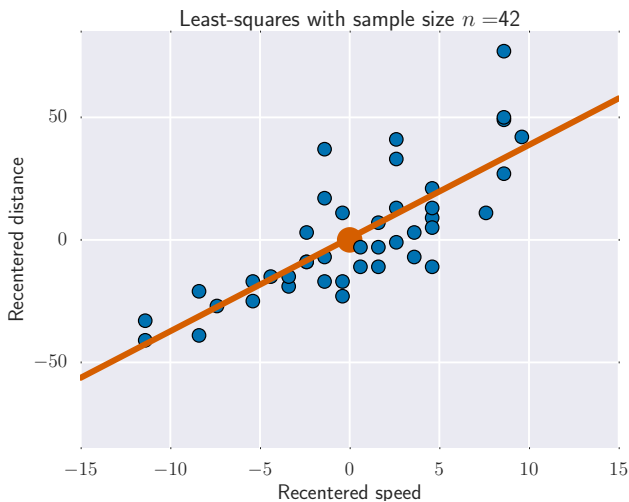
Extreme points – leverage effect (II)



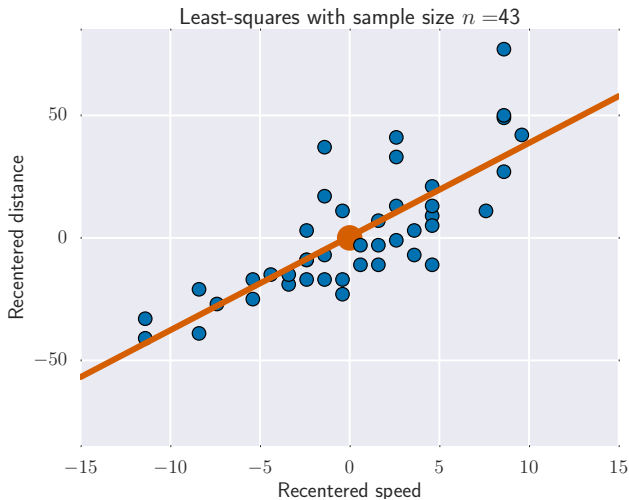
Extreme points – leverage effect (II)



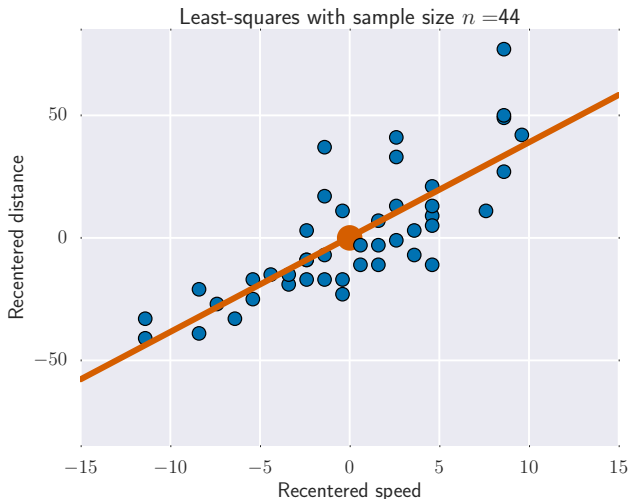
Extreme points – leverage effect (II)



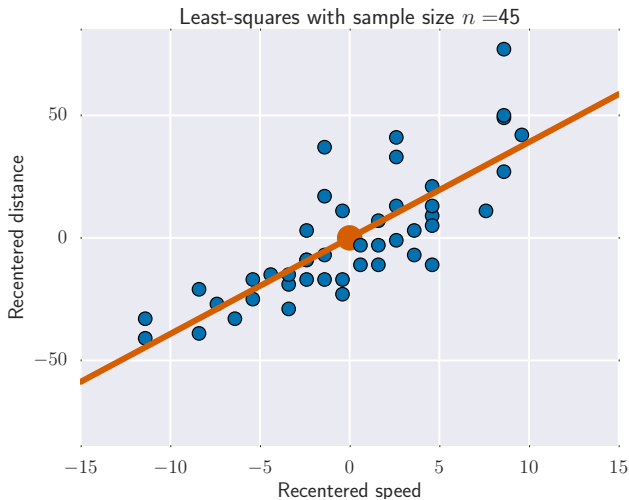
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



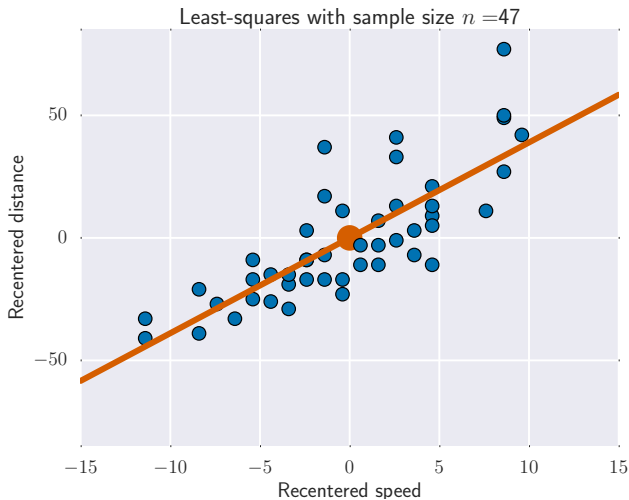
Extreme points – leverage effect (II)



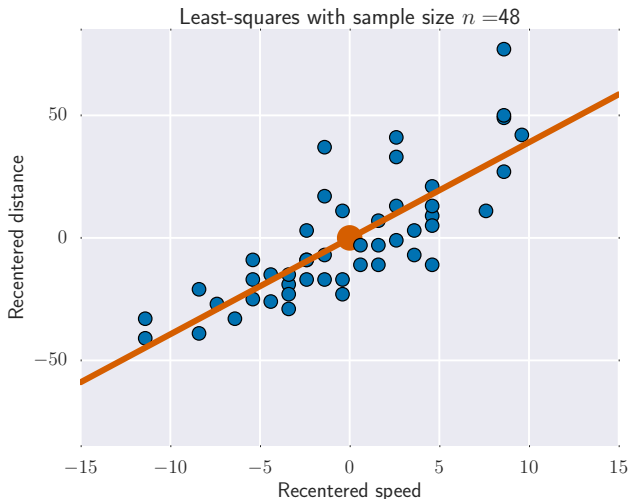
Extreme points – leverage effect (II)



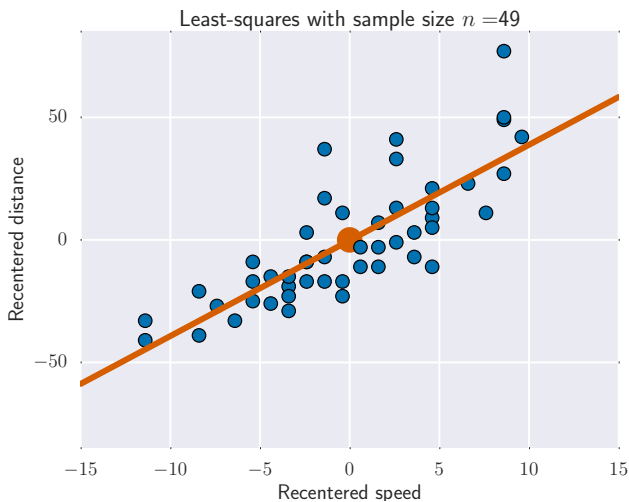
Extreme points – leverage effect (II)



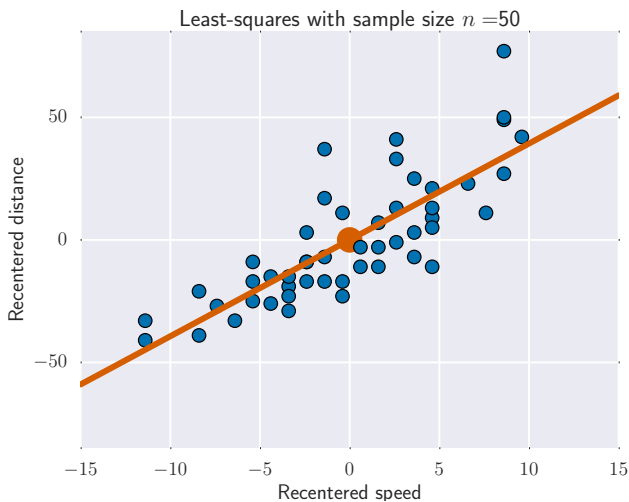
Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



Extreme points – leverage effect (II)



Centering + scaling (standardization)

Centered-scaled model :

$$\forall i = 1, \dots, n : \begin{cases} x_i'' = (x_i - \bar{x}_n) / \sqrt{\text{var}_n(\mathbf{x})} \\ y_i'' = (y_i - \bar{y}_n) / \sqrt{\text{var}_n(\mathbf{y})} \end{cases} \Leftrightarrow \begin{cases} \mathbf{x}'' = \frac{\mathbf{x} - \bar{x}_n \mathbf{1}_n}{\sqrt{\text{var}_n(\mathbf{x})}} \\ \mathbf{y}'' = \frac{\mathbf{y} - \bar{y}_n \mathbf{1}_n}{\sqrt{\text{var}_n(\mathbf{y})}} \end{cases}$$

Solving OLS with $(\mathbf{x}'', \mathbf{y}'')$ then

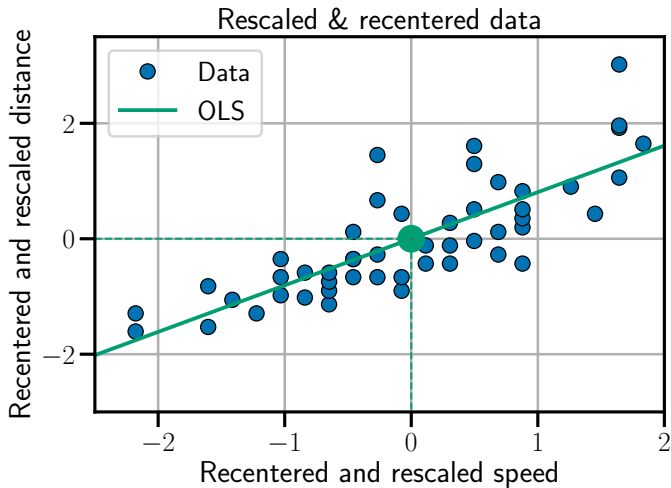
$$\begin{cases} \hat{\theta}_0'' = 0 \\ \hat{\theta}_1'' = \frac{1}{n} \sum_{i=1}^n x_i'' y_i'' \end{cases}$$

Rem: equivalent to choosing the points cloud center of mass as origin and normalize \mathbf{x} and \mathbf{y} to have unit **empirical norm** $\|\cdot\|_n$:

$$\|\mathbf{x}''\|_n^2 = \frac{1}{n} \sum_{i=1}^n (x_i'')^2 = 1$$


$$\|\mathbf{y}''\|_n^2 = \frac{1}{n} \sum_{i=1}^n (y_i'')^2 = 1$$

Centering + scaling



When/why preprocessing ?

Centering y or using an intercept (or adding a constant feature) is equivalent

Rem: for sparse ( : *creux*) cases centering y adding a constant feature could be preferred

Scaling features is important :

- ▶ if you want to interpret the coefficients' amplitude in regression (better solution : t-tests)
- ▶ if you want to penalize or regularize coefficients (*cf.* Lasso, Ridge, etc.) a single scale is needed
- ▶ for computing reasons (e.g., store scaling to improve efficiency, etc.)

Rem: in practice centering/scaling is useful for **estimation** not so much for **prediction** (see next courses)

What happens with the logarithm scaling ?

Centering with Python

Use centering classes from sklearn, see preprocessing :
<http://scikit-learn.org/stable/modules/preprocessing.html>

```
from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X)

print(np.isclose(scaler.mean_, np.mean(X)))

print(np.array_equal(scaler.std_, np.std(X)))

print(np.array_equal(scaler.transform(X),
                     (X - np.mean(X)) / np.std(X)))

print(np.array_equal(scaler.transform([26]),
                     (26 - np.mean(X)) / np.std(X)))
```

Rem: most valuable with pipeline

<http://scikit-learn.org/stable/modules/pipeline.html>

Definitions

We call **prediction** function the function that associates an estimation of the variable of interest to a new sample. For least squares the prediction is given by :

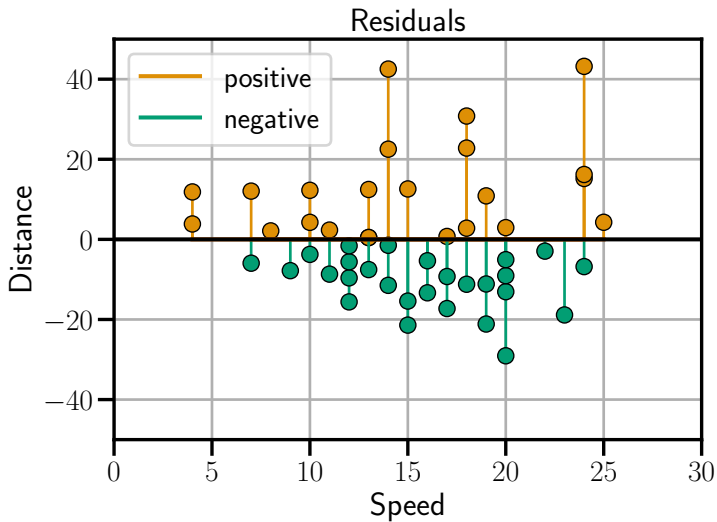
$$\text{pred}(x_{n+1}) = \hat{\theta}_0 + \hat{\theta}_1 x_{n+1}$$

Rem: often written \hat{y}_{n+1} (implicit dependence on x_{n+1}) The **residual** : difference between observations and predicted values

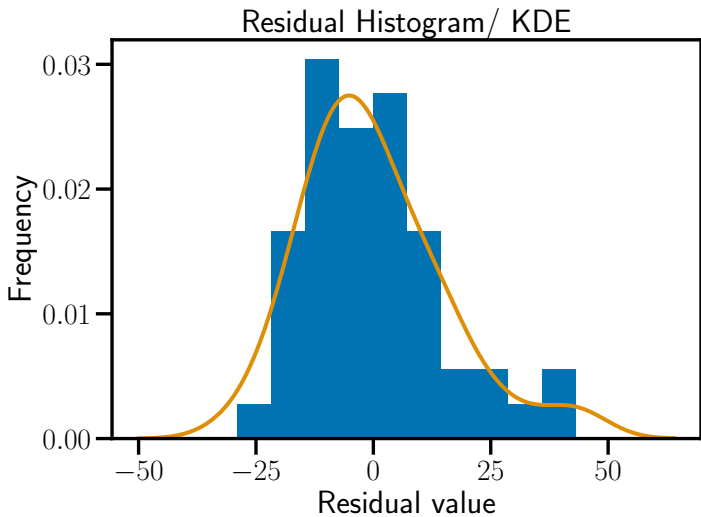
$$r_i = y_i - \text{pred}(x_i) = y_i - \hat{y}_i = y_i - (\hat{\theta}_0 + \hat{\theta}_1 x_i)$$

Rem: observable estimate of the unobservable statistical error

Residuals (on cars)



Residual histograms



Least squares motivation

- ▶ Computing advantage : computationally heavy methods avoided before computers (e.g., iterative methods)
- ▶ Theoretical advantage : least square analysis easy under simple hypothesis
- ▶ Interpretability : how much does the regressor increase with the features

Example : under additive white Gaussian noise assumption *i.e.*, $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ the maximum likelihood is equivalent to solving least squares to estimate (θ_0^*, θ_1^*)

Rem: for another noise model and/or to limit outliers influence one can solve (see e.g., QuantReg in statsmodels)

$$\hat{\theta} = (\hat{\theta}_0, \hat{\theta}_1) \in \arg \min_{(\theta_0, \theta_1) \in \mathbb{R}^2} \sum_{i=1}^n |y_i - \theta_0 - \theta_1 x_i|$$

Gaussian likelihood

Rem: univariate probability density function (pdf) We write $Y \sim \mathcal{N}(\mu, \sigma^2)$, for a random variable with pdf

$$\varphi_{\mu, \sigma}(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

Assume : $y_i \sim \mathcal{N}(\theta_0^* + \theta_1^* x_i, \sigma^2)$, i.e., $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$, then the most **likely** couple (θ_0, θ_1) based on the observations is maximizing the pdf of (y_1, \dots, y_n)

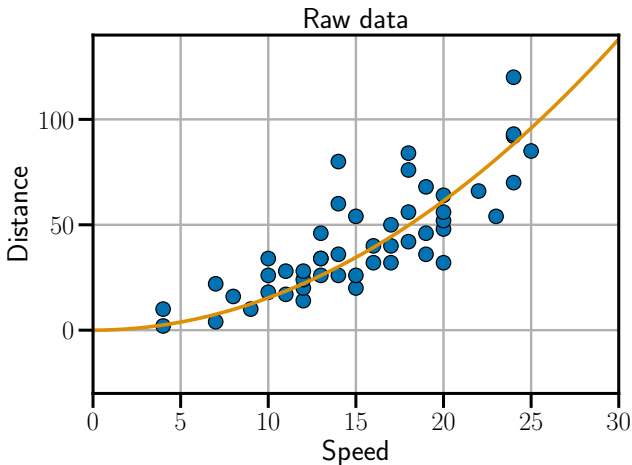
Under an independence hypothesis, this is achieved by solving :

$$(\hat{\theta}_0, \hat{\theta}_1) \in \arg \max_{(\theta_0, \theta_1) \in \mathbb{R}^2} \prod_{i=1}^n \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \theta_0 - \theta_1 x_i)^2}{2\sigma^2}\right) \right)$$

Exercise: check that this is equivalent to the least squares formulation

Discussion : toward multivariate cases

Physical laws (or your driving school memories) would lead to rather pick a **quadratic** model instead of a **linear** one : the OLS can be applied by choosing x_i^2 as features instead of x_i :



Web sites and books to go further

- ▶ Datascience in general : Blog + videos by Jake Vanderplas
<http://jakevdp.github.io/>
Homework for next lesson : watch the following videos
<http://jakevdp.github.io/blog/2017/03/03/reproducible-data-analysis-in-jupyter/>
- ▶ A few [notebooks](#) of OLS with statsmodels
- ▶ [McKinney \(2012\)](#) about Python for statistics
- ▶ [Lejeune \(2010\)](#) about linear models (in French)
- ▶ Regression course by [B. Delyon](#) (in French, more technical)

References I

- ▶ D. P. Bertsekas.
Nonlinear programming.
Athena Scientific, 1999.
- ▶ S. Boyd and L. Vandenberghe.
Convex optimization.
Cambridge University Press, Cambridge, 2004.
- ▶ B. Delyon.
Régression, 2015.
[https://perso.univ-rennes1.fr/bernard.delyon/
regression.pdf](https://perso.univ-rennes1.fr/bernard.delyon/regression.pdf).
- ▶ D. Foata and A. Fuchs.
Calcul des probabilités : cours et exercices corrigés.
Masson, 1996.

References II

- ▶ G. H. Golub and C. F. van Loan.
Matrix computations.
Johns Hopkins University Press, Baltimore, MD, fourth edition, 2013.
- ▶ R. A. Horn and C. R. Johnson.
Topics in matrix analysis.
Cambridge University Press, Cambridge, 1994.
Corrected reprint of the 1991 original.
- ▶ M. Lejeune.
Statistiques, la théorie et ses applications.
Springer, 2010.
- ▶ W. McKinney.
Python for Data Analysis : Data Wrangling with Pandas, NumPy, and IPython.
O'Reilly Media, 2012.

References III

- ▶ K. P. Murphy.
Machine learning : a probabilistic perspective.
MIT press, 2012.